

more pandas

Lecture 09

Dr. Colin Rundel

Index objects

Columns and Indexes

When constructing a DataFrame we can specify the indexes for both the rows (`index`) and columns (`columns`),

```
1 df = pd.DataFrame(  
2     np.random.randn(5, 3),  
3     columns=['A', 'B', 'C']  
4 )  
5 df
```

	A	B	C
0	1.214010	1.293817	-2.529983
1	-0.979902	-0.985834	-1.479055
2	-0.458005	0.527957	-0.402826
3	1.488340	0.380165	0.020759
4	-1.092834	-0.125232	0.625418

```
1 df.columns
```

```
Index(['A', 'B', 'C'], dtype='object')
```

```
1 df.index
```

```
RangeIndex(start=0, stop=5, step=1)
```

```
1 df = pd.DataFrame(  
2     np.random.randn(3, 3),  
3     index=['x', 'y', 'z'],  
4     columns=['A', 'B', 'C']  
5 )  
6 df
```

	A	B	C
x	0.777821	2.477700	0.814513
y	1.737038	-1.112535	-0.723814
z	0.041552	-1.455307	0.417292

```
1 df.columns
```

```
Index(['A', 'B', 'C'], dtype='object')
```

```
1 df.index
```

```
Index(['x', 'y', 'z'], dtype='object')
```

Index objects

pandas' `Index` class and its subclasses provide the infrastructure necessary for lookups, data alignment, and other related tasks. You can think of them as being an immutable *multiset* (duplicate values are allowed).

```
1 pd.Index(['A', 'B', 'C'])
```

```
Index(['A', 'B', 'C'], dtype='object')
```

```
1 pd.Index(['A', 'B', 'C', 'A'])
```

```
Index(['A', 'B', 'C', 'A'], dtype='object')
```

```
1 pd.Index(range(5))
```

```
RangeIndex(start=0, stop=5, step=1)
```

```
1 pd.Index(list(range(5)))
```

```
Int64Index([0, 1, 2, 3, 4], dtype='int64')
```

Indexes as sets

While it is not something you will need to do very often, since Indexes are “sets” the various set operations and methods are available.

```
1 a = pd.Index(['c', 'b', 'a'])
2 b = pd.Index(['c', 'e', 'd'])
```

```
1 a.union(b)
```

```
Index(['a', 'b', 'c', 'd', 'e'], dtype='object')
```

```
1 a.intersection(b)
```

```
Index(['c'], dtype='object')
```

```
1 c = pd.Index([1.0, 1.5, 2.0])
2 d = pd.Index(range(5))
3
4 c.union(d)
```

```
Float64Index([0.0, 1.0, 1.5, 2.0, 3.0, 4.0],
dtype='float64')
```

```
1 a.difference(b)
```

```
Index(['a', 'b'], dtype='object')
```

```
1 a.symmetric_difference(b)
```

```
Index(['a', 'b', 'd', 'e'], dtype='object')
```

```
1 e = pd.Index(["A", "B", "C"])
2 f = pd.Index(range(5))
3
4 e.union(f)
```

```
Index(['A', 'B', 'C', 0, 1, 2, 3, 4],
dtype='object')
```

Index metadata

You can attach names to an index, which will then show when displaying the DataFrame or Index,

```
1 df = pd.DataFrame(  
2     np.random.randn(3, 3),  
3     index=pd.Index(['x', 'y', 'z'], name="rows"),  
4     columns=pd.Index(['A', 'B', 'C'], name="cols")  
5 )  
6 df
```

cols	A	B	C
rows			
x	-0.779489	1.053635	1.439538
y	1.732166	-0.471819	0.437045
z	-1.624165	-0.074081	1.694972

```
1 df.columns
```

```
Index(['A', 'B', 'C'], dtype='object', name='cols')
```

```
1 df.index
```

```
Index(['x', 'y', 'z'], dtype='object', name='rows')
```

```
1 df.columns.rename("m")
```

```
Index(['A', 'B', 'C'], dtype='object', name='m')
```

```
1 df.index.set_names("n")
```

```
Index(['x', 'y', 'z'], dtype='object', name='n')
```

```
1 df
```

cols	A	B	C
rows			
x	-0.779489	1.053635	1.439538
y	1.732166	-0.471819	0.437045
z	-1.624165	-0.074081	1.694972

Renaming indexes inplace

If you want to change the index names inplace either assign directly to the `name` attribute or use the `inplace=True` argument with `rename()`.

```
1 df
```

cols	A	B	C
x	-0.779489	1.053635	1.439538
y	1.732166	-0.471819	0.437045
z	-1.624165	-0.074081	1.694972

```
1 df.columns.name = "o"  
2 df.index.name = "p"  
3 df
```

o	A	B	C
p			
x	-0.779489	1.053635	1.439538
y	1.732166	-0.471819	0.437045
z	-1.624165	-0.074081	1.694972

```
1 df.columns.rename("q", inplace=True)  
2 df.index.rename("r", inplace=True)  
3 df
```

q	A	B	C
r			
x	-0.779489	1.053635	1.439538
y	1.732166	-0.471819	0.437045
z	-1.624165	-0.074081	1.694972

Indexes and missing values

It is possible for an index to contain missing values (e.g. `np.nan`) but this is generally a bad idea and should be avoided.

```
1 pd.Index([1,2,3,np.nan,5])
```

```
Float64Index([1.0, 2.0, 3.0, nan, 5.0], dtype='float64')
```

```
1 pd.Index(["A","B",np.nan,"D", None])
```

```
Index(['A', 'B', nan, 'D', None], dtype='object')
```

Missing values can be replaced via the `fillna()` method,

```
1 pd.Index([1,2,3,np.nan,5]).fillna(0)
```

```
Float64Index([1.0, 2.0, 3.0, 0.0, 5.0], dtype='float64')
```

```
1 pd.Index(["A","B",np.nan,"D", None]).fillna("Z")
```

```
Index(['A', 'B', 'Z', 'D', 'Z'], dtype='object')
```


Changing a DataFrame's index

Existing columns can be made an index via `set_index()` and removed via `reset_index()`,

```
1 data
```

	a	b	c	d
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3
3	foo	two	w	4

```
1 data.set_index('a')
```

	b	c	d
a			
bar	one	z	1
bar	two	y	2
foo	one	x	3
foo	two	w	4

```
1 data.set_index('c', drop=False)
```

	a	b	c	d
c				
z	bar	one	z	1
y	bar	two	y	2
x	foo	one	x	3
w	foo	two	w	4

```
1 data.set_index('a').reset_index()
```

	a	b	c	d
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3
3	foo	two	w	4

```
1 data.set_index('c').reset_index(drop=True)
```

	a	b	d
0	bar	one	1
1	bar	two	2
2	foo	one	3
3	foo	two	4

Creating a new index

New index values can be attached to a DataFrame via `reindex()`,

```
1 data
```

	a	b	c	d
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3
3	foo	two	w	4

```
1 data.reindex(["w","x","y","z"])
```

	a	b	c	d
w	NaN	NaN	NaN	NaN
x	NaN	NaN	NaN	NaN
y	NaN	NaN	NaN	NaN
z	NaN	NaN	NaN	NaN

```
1 data.reindex(range(5,-1,-1))
```

	a	b	c	d
5	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN
3	foo	two	w	4.0
2	foo	one	x	3.0
1	bar	two	y	2.0
0	bar	one	z	1.0

```
1 data.reindex(columns = ["a","b","c","d","e"])
```

	a	b	c	d	e
0	bar	one	z	1	NaN
1	bar	two	y	2	NaN
2	foo	one	x	3	NaN
3	foo	two	w	4	NaN

```
1 data.index = ["w","x","y","z"]  
2 data
```

	a	b	c	d
w	bar	one	z	1
x	bar	two	y	2
y	foo	one	x	3
z	foo	two	w	4

Renaming levels

Alternatively, row or column index levels can be renamed via `rename()`,

```
1 data
```

	a	b	c	d
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3
3	foo	two	w	4

```
1 data.rename(index = pd.Series(["m","n","o","p"]))
```

	a	b	c	d
m	bar	one	z	1
n	bar	two	y	2
o	foo	one	x	3
p	foo	two	w	4

```
1 data.rename_axis(index="rows")
```

	a	b	c	d
rows				
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3
3	foo	two	w	4

```
1 data.rename(columns = {"a":"w", "b":"x",  
2                       "c":"y", "d":"z"})
```

	w	x	y	z
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3
3	foo	two	w	4

```
1 data.rename_axis(columns="cols")
```

	a	b	c	d
cols				
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3
3	foo	two	w	4

MultiIndexes

MultiIndex objects

These are a hierarchical analog of standard Index objects, there are a number of methods for constructing them based on the initial object

```
1 tuples = [('A','x'), ('A','y'),
2           ('B','x'), ('B','y'),
3           ('C','x'), ('C','y')]
4 pd.MultiIndex.from_tuples(
5     tuples, names=["1st","2nd"]
6 )
```

```
MultiIndex([('A', 'x'),
            ('A', 'y'),
            ('B', 'x'),
            ('B', 'y'),
            ('C', 'x'),
            ('C', 'y')],
            names=['1st', '2nd'])
```

```
1 pd.MultiIndex.from_product(
2     [ ["A","B","C"], ["x","y"] ], names=["1st","2nd"]
3 )
```

```
MultiIndex([('A', 'x'),
            ('A', 'y'),
            ('B', 'x'),
            ('B', 'y'),
            ('C', 'x'),
            ('C', 'y')],
            names=['1st', '2nd'])
```

DataFrame with MultiIndex

```
1 idx = pd.MultiIndex.from_tuples(  
2     tuples, names=["1st", "2nd"]  
3 )  
4  
5 pd.DataFrame(  
6     np.random.rand(6, 2),  
7     index = idx,  
8     columns=["m", "n"]  
9 )
```

		m	n
1st	2nd		
A	x	0.477958	0.110635
	y	0.176785	0.524042
B	x	0.686342	0.254666
	y	0.366431	0.335209
C	x	0.526298	0.868879
	y	0.281066	0.523480

Column MultiIndex

MultiIndexes can also be used for columns (or both rows and columns),

```
1 cidx = pd.MultiIndex.from_product(  
2     [ ["A","B"], ["x","y"] ], names=["c1","c2"]  
3 )  
4  
5 pd.DataFrame(  
6     np.random.rand(4,4), columns = cidx  
7 )
```

c1	A		B	
c2	x	y	x	y
0	0.468450	0.956413	0.703818	0.262485
1	0.628330	0.063963	0.837249	0.867690
2	0.280684	0.258858	0.810500	0.730225
3	0.590576	0.607856	0.573620	0.788613

```
1 ridx = pd.MultiIndex.from_product(  
2     [ ["m","n"], ["l","p"] ], names=["r1","r2"]  
3 )  
4  
5 pd.DataFrame(  
6     np.random.rand(4,4),  
7     index= ridx, columns = cidx  
8 )
```

c1	A		B	
c2	x	y	x	y
r1 r2				
m l	0.503954	0.687235	0.571731	0.062033
p	0.890484	0.855219	0.250289	0.254360
n l	0.131349	0.445940	0.408909	0.366915
p	0.589651	0.551470	0.995567	0.952875

MultiIndex indexing

```
1 data
```

c1		A		B	
c2		x	y	x	y
r1	r2				
m	l	0.558355	0.802381	0.501756	0.213913
	p	0.783600	0.142549	0.129497	0.972312
n	l	0.976518	0.737849	0.054503	0.756322
	p	0.429930	0.924574	0.291695	0.494732

```
1 data["A"]
```

c2		x	y
r1	r2		
m	l	0.558355	0.802381
	p	0.783600	0.142549
n	l	0.976518	0.737849
	p	0.429930	0.924574

```
1 data["x"]
```

Error: KeyError: 'x'

```
1 data["m"]
```

Error: KeyError: 'm'

```
1 data["m","A"]
```

Error: KeyError: ('m', 'A')

```
1 data["A","x"]
```

r1	r2	
m	l	0.558355
	p	0.783600
n	l	0.976518
	p	0.429930

Name: (A, x), dtype: float64

```
1 data["A"]["x"]
```

r1	r2	
m	l	0.558355
	p	0.783600
n	l	0.976518
	p	0.429930

Name: x, dtype: float64

MultiIndex indexing via `iloc`

```
1 data.iloc[0]
```

```
c1  c2
A   x    0.558355
    y    0.802381
B   x    0.501756
    y    0.213913
Name: (m, l), dtype: float64
```

```
1 data.iloc[(0,1)]
```

```
0.8023809565660679
```

```
1 data.iloc[[0,1]]
```

```
c1          A          B
c2          x          y          x          y
r1 r2
m  l  0.558355  0.802381  0.501756  0.213913
   p  0.783600  0.142549  0.129497  0.972312
```

```
1 data.iloc[:,0]
```

```
r1  r2
m   l    0.558355
    p    0.783600
n   l    0.976518
    p    0.429930
Name: (A, x), dtype: float64
```

```
1 data.iloc[0,1]
```

```
0.8023809565660679
```

```
1 data.iloc[0,[0,1]]
```

```
c1  c2
A   x    0.558355
    y    0.802381
Name: (m, l), dtype: float64
```

MultiIndex indexing via `loc`

```
1 data.loc["m"]
```

c1	A		B	
c2	x	y	x	y
l	0.558355	0.802381	0.501756	0.213913
p	0.783600	0.142549	0.129497	0.972312

```
1 data.loc["l"]
```

Error: KeyError: 'l'

```
1 data.loc[:, "A"]
```

c2		x	y
r1	r2		
m	l	0.558355	0.802381
	p	0.783600	0.142549
n	l	0.976518	0.737849
	p	0.429930	0.924574

```
1 data.loc[("m", "l")]
```

c1	c2	
A	x	0.558355
	y	0.802381
B	x	0.501756
	y	0.213913

Name: (m, l), dtype: float64

```
1 data.loc[:, ("A", "y")]
```

r1	r2	
m	l	0.802381
	p	0.142549
n	l	0.737849
	p	0.924574

Name: (A, y), dtype: float64

Fancier indexing with `loc`

Index slices can also be used with combinations of indexes and index tuples,

```
1 data.loc["m":"n"]
```

c1		A		B	
c2		x	y	x	y
r1	r2				
m	l	0.558355	0.802381	0.501756	0.213913
	p	0.783600	0.142549	0.129497	0.972312
n	l	0.976518	0.737849	0.054503	0.756322
	p	0.429930	0.924574	0.291695	0.494732

```
1 data.loc[("m","l"):(("n","l"))]
```

c1		A		B	
c2		x	y	x	y
r1	r2				
m	l	0.558355	0.802381	0.501756	0.213913
	p	0.783600	0.142549	0.129497	0.972312
n	l	0.976518	0.737849	0.054503	0.756322

```
1 data.loc[("m","p"):"n"]
```

c1		A		B	
c2		x	y	x	y
r1	r2				
m	p	0.783600	0.142549	0.129497	0.972312
n	l	0.976518	0.737849	0.054503	0.756322
	p	0.429930	0.924574	0.291695	0.494732

```
1 data.loc[[("m","p"),("n","l")]]
```

c1		A		B	
c2		x	y	x	y
r1	r2				
m	p	0.783600	0.142549	0.129497	0.972312
n	l	0.976518	0.737849	0.054503	0.756322

Selecting nested levels

The previous methods don't give easy access to indexing on nested index levels, this is possible via the cross-section method `xs()`,

```
1 data.xs("p", level="r2")
```

c1	A		B	
c2	x	y	x	y
r1				
m	0.78360	0.142549	0.129497	0.972312
n	0.42993	0.924574	0.291695	0.494732

```
1 data.xs("m", level="r1")
```

c1	A		B	
c2	x	y	x	y
r2				
l	0.558355	0.802381	0.501756	0.213913
p	0.783600	0.142549	0.129497	0.972312

```
1 data.xs("y", level="c2", axis=1)
```

c1	A		B	
r1	r2			
m	l	0.802381	0.213913	
	p	0.142549	0.972312	
n	l	0.737849	0.756322	
	p	0.924574	0.494732	

```
1 data.xs("B", level="c1", axis=1)
```

c2	x		y	
r1	r2			
m	l	0.501756	0.213913	
	p	0.129497	0.972312	
n	l	0.054503	0.756322	
	p	0.291695	0.494732	

Setting MultiIndexes

It is also possible to construct a MultiIndex or modify an existing one using `set_index()` and `reset_index()`,

```
1 data
```

	a	b	c	d
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3

```
1 data.set_index(['a','b'])
```

		c	d
a	b		
bar	one	z	1
	two	y	2
foo	one	x	3

```
1 data.set_index('c', append=True)
```

	a	b	d	
c				
0	z	bar	one	1
1	y	bar	two	2
2	x	foo	one	3

```
1 data.set_index(['a','b']).reset_index()
```

	a	b	c	d
0	bar	one	z	1
1	bar	two	y	2
2	foo	one	x	3

```
1 data.set_index(['a','b']).reset_index(level=1)
```

		b	c	d
a				
bar	one	z	1	
bar	two	y	2	
foo	one	x	3	

Reshaping data

Long to wide (pivot)

```
1 df
```

	country	year	type	count
0	A	1999	cases	0.7K
1	A	1999	pop	19M
2	A	2000	cases	2K
3	A	2000	pop	20M
4	B	1999	cases	37K
5	B	1999	pop	172M
6	B	2000	cases	80K
7	B	2000	pop	174M
8	C	1999	cases	212K
9	C	1999	pop	1T
10	C	2000	cases	213K
11	C	2000	pop	1T

```
1 df_wide = df.pivot(  
2     index=["country","year"],  
3     columns="type",  
4     values="count"  
5 )  
6 df_wide
```

		type	cases	pop
	country	year		
A		1999	0.7K	19M
		2000	2K	20M
B		1999	37K	172M
		2000	80K	174M
C		1999	212K	1T
		2000	213K	1T

pivot indexes

```
1 df_wide.index
```

```
MultiIndex([( 'A', 1999),  
            ( 'A', 2000),  
            ( 'B', 1999),  
            ( 'B', 2000),  
            ( 'C', 1999),  
            ( 'C', 2000)],  
           names=[ 'country', 'year'])
```

```
1 df_wide.columns
```

```
Index([ 'cases', 'pop'], dtype='object',  
      name='type')
```

```
1 df_wide.reset_index(  
2 ).rename_axis(  
3     columns=None  
4 )
```

	country	year	cases	pop
0	A	1999	0.7K	19M
1	A	2000	2K	20M
2	B	1999	37K	172M
3	B	2000	80K	174M
4	C	1999	212K	1T
5	C	2000	213K	1T

Wide to long (melt)

```
1 df
```

	country	1999	2000
0	A	0.7K	2K
1	B	37K	80K
2	C	212K	213K

```
1 df_long = df.melt(  
2     id_vars="country",  
3     var_name="year"  
4 )  
5 df_long
```

	country	year	value
0	A	1999	0.7K
1	B	1999	37K
2	C	1999	212K
3	A	2000	2K
4	B	2000	80K
5	C	2000	213K

Separate Example - splits and explosions

```
1 df
```

	country	year	rate
0	A	1999	0.7K/19M
1	A	2000	2K/20M
2	B	1999	37K/172M
3	B	2000	80K/174M
4	C	1999	212K/1T
5	C	2000	213K/1T

```
1 df.assign(  
2     rate = lambda d: d.rate.str.split("/")  
3 )
```

	country	year	rate
0	A	1999	[0.7K, 19M]
1	A	2000	[2K, 20M]
2	B	1999	[37K, 172M]
3	B	2000	[80K, 174M]
4	C	1999	[212K, 1T]
5	C	2000	[213K, 1T]

```

1 ( df.assign(
2     rate = lambda d: d.rate.str.split("/")
3 )
4 .explode("rate")
5 .assign(
6     type = lambda d: ["cases", "pop"] * int(d.sh
7 )
8 )

```

	country	year	rate	type
0	A	1999	0.7K	cases
0	A	1999	19M	pop
1	A	2000	2K	cases
1	A	2000	20M	pop
2	B	1999	37K	cases
2	B	1999	172M	pop
3	B	2000	80K	cases
3	B	2000	174M	pop
4	C	1999	212K	cases
4	C	1999	1T	pop
5	C	2000	213K	cases
5	C	2000	1T	pop

Putting it together

```
1 ( df
2   .assign(
3     rate = lambda d: d.rate.str.split("/")
4   )
5   .explode("rate")
6   .assign(
7     type = lambda d: ["cases", "pop"] *
8                     int(d.shape[0]/2)
9   )
10 )
```

	country	year	rate	type
0	A	1999	0.7K	cases
0	A	1999	19M	pop
1	A	2000	2K	cases
1	A	2000	20M	pop
2	B	1999	37K	cases
2	B	1999	172M	pop
3	B	2000	80K	cases
3	B	2000	174M	pop
4	C	1999	212K	cases
4	C	1999	1T	pop
5	C	2000	213K	cases
5	C	2000	1T	pop

```
1 ( df.assign(
2     rate = lambda d: d.rate.str.split("/")
3   )
4   .explode("rate")
5   .assign(
6     type = lambda d: ["cases", "pop"] *
7                     int(d.shape[0]/2)
8   )
9   .pivot(
10    index=["country", "year"],
11    columns="type",
12    values="rate"
13  )
14  .reset_index()
15 )
```

	type	country	year	cases	pop
0		A	1999	0.7K	19M
1		A	2000	2K	20M
2		B	1999	37K	172M
3		B	2000	80K	174M
4		C	1999	212K	1T
5		C	2000	213K	1T

Separate Example - A better way

```
1 df
```

	country	year	rate
0	A	1999	0.7K/19M
1	A	2000	2K/20M
2	B	1999	37K/172M
3	B	2000	80K/174M
4	C	1999	212K/1T
5	C	2000	213K/1T

```
1 df.assign(  
2     counts = lambda d: d.rate.str.split("/").str[0]  
3     pop     = lambda d: d.rate.str.split("/").str[1]  
4 )
```

	country	year	rate	counts	pop
0	A	1999	0.7K/19M	0.7K	19M
1	A	2000	2K/20M	2K	20M
2	B	1999	37K/172M	37K	172M
3	B	2000	80K/174M	80K	174M
4	C	1999	212K/1T	212K	1T
5	C	2000	213K/1T	213K	1T

If you dont want to repeat the split,

```
1 df.assign(  
2     rate = lambda d: d.rate.str.split("/"),  
3     counts = lambda d: d.rate.str[0],  
4     pop     = lambda d: d.rate.str[1]  
5 ).drop("rate", axis=1)
```

	country	year	counts	pop
0	A	1999	0.7K	19M
1	A	2000	2K	20M
2	B	1999	37K	172M
3	B	2000	80K	174M
4	C	1999	212K	1T
5	C	2000	213K	1T

Exercise 1

Create a DataFrame from the data available at https://sta663-sp23.github.io/slides/data/us_rent.csv using `pd.read_csv()`.

These data come from the 2017 American Community Survey and reflect the following values:

- * `name` - name of state
- * `variable` - Variable name: income = median yearly income, rent = median monthly rent
- * `estimate` - Estimated value
- * `moe` - 90% margin of error

Using these data find the state(s) with the lowest income to rent ratio.

Split-Apply-Combine

groupby

Groups can be created within a DataFrame via `groupby()` - these groups are then used by the standard summary methods (e.g. `sum()`, `mean()`, `std()`, etc.).

```
1 cereal = pd.read_csv("https://sta663-sp23.github.io/slides/data/cereal.csv")
2 cereal
```

	name	mfr	...	sugars	rating
0	100% Bran	Nabisco	...	6	68.402973
1	100% Natural Bran	Quaker Oats	...	8	33.983679
2	All-Bran	Kellogg's	...	5	59.425505
3	All-Bran with Extra Fiber	Kellogg's	...	0	93.704912
4	Almond Delight	Ralston Purina	...	8	34.384843
..
72	Triples	General Mills	...	3	39.106174
73	Trix	General Mills	...	12	27.753301
74	Wheat Chex	Ralston Purina	...	3	49.787445
75	Wheaties	General Mills	...	3	51.592193
76	Wheaties Honey Gold	General Mills	...	8	36.187559

[77 rows x 6 columns]

```
1 cereal.groupby("type")
```

<pandas.core.groupby.generic.DataFrameGroupBy object at 0x290553850>

GroupBy attributes and methods

```
1 cereal.groupby("type").groups
```

```
{'Cold': [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76], 'Hot': [20, 43, 57]}
```

```
1 cereal.groupby("type").mean(numeric_only=True)
```

	calories	sugars	rating
type			
Cold	107.162162	7.175676	42.095218
Hot	100.000000	1.333333	56.737708

```
1 cereal.groupby("mfr").groups
```

```
{'General Mills': [5, 7, 11, 12, 13,
14, 18, 22, 31, 36, 40, 42, 47, 51, 59,
69, 70, 71, 72, 73, 75, 76],
'Kellogg's': [2, 3, 6, 16, 17, 19, 21,
24, 25, 26, 28, 38, 39, 46, 48, 49, 50,
53, 58, 60, 62, 66, 67], 'Maltex':
[43], 'Nabisco': [0, 20, 63, 64, 65,
68], 'Post': [9, 27, 29, 30, 32, 33,
34, 37, 52], 'Quaker Oats': [1, 10, 35,
41, 54, 55, 56, 57], 'Ralston Purina':
[4, 8, 15, 23, 44, 45, 61, 74]}
```

```
1 cereal.groupby("mfr").size()
```

```
mfr
General Mills      22
Kellogg's          23
Maltex             1
Nabisco            6
Post              9
Quaker Oats        8
Ralston Purina     8
dtype: int64
```

Selecting groups

Groups can be accessed via `get_group()` or the `DataFrameGroupBy` can be iterated over,

```
1 cereal.groupby("type").get_group("Hot")
```

	name	mfr	type	calories	sugars	rating
20	Cream of Wheat (Quick)	Nabisco	Hot	100	0	64.533816
43	Maypo	Maltex	Hot	100	3	54.850917
57	Quaker Oatmeal	Quaker Oats	Hot	100	1	50.828392

```
1 cereal.groupby("mfr").get_group("Post")
```

	name	mfr	...	sugars	rating
9	Bran Flakes	Post	...	5	53.313813
27	Fruit & Fibre Dates; Walnuts; and Oats	Post	...	10	40.917047
29	Fruity Pebbles	Post	...	12	28.025765
30	Golden Crisp	Post	...	15	35.252444
32	Grape Nuts Flakes	Post	...	5	52.076897
33	Grape-Nuts	Post	...	3	53.371007
34	Great Grains Pecan	Post	...	4	45.811716
37	Honey-comb	Post	...	11	28.742414
52	Post Nat. Raisin Bran	Post	...	14	37.840594

[9 rows x 6 columns]

Iterating groups

```
1 for name, group in cereal.groupby("type"):
2     print(f"# {name}\n{group}\n\n")
```

Cold

	name	mfr	...	sugars	rating
0	100% Bran	Nabisco	...	6	68.402973
1	100% Natural Bran	Quaker Oats	...	8	33.983679
2	All-Bran	Kellogg's	...	5	59.425505
3	All-Bran with Extra Fiber	Kellogg's	...	0	93.704912
4	Almond Delight	Ralston Purina	...	8	34.384843
..
72	Triples	General Mills	...	3	39.106174
73	Trix	General Mills	...	12	27.753301
74	Wheat Chex	Ralston Purina	...	3	49.787445
75	Wheaties	General Mills	...	3	51.592193
76	Wheaties Honey Gold	General Mills	...	8	36.187559

[74 rows x 6 columns]

Aggregation

The `aggregate()` function or `agg()` method can be used to compute summary statistics for each group,

```
1 cereal.groupby("mfr").agg("mean")
```

	calories	sugars	rating
mfr			
General Mills	111.363636	7.954545	34.485852
Kellogg's	108.695652	7.565217	44.038462
Maltex	100.000000	3.000000	54.850917
Nabisco	86.666667	1.833333	67.968567
Post	108.888889	8.777778	41.705744
Quaker Oats	95.000000	5.500000	42.915990
Ralston Purina	115.000000	6.125000	41.542997

<string>:1: FutureWarning: The default value of `numeric_only` in `DataFrameGroupBy.mean` is deprecated. In a future version, `numeric_only` will default to `False`. Either specify `numeric_only` or select only columns which should be valid for the function.

Aggregation with multiple functions

```
1 cereal.groupby("mfr").agg([np.mean, np.std])
```

	calories		sugars		rating	
	mean	std	mean	std	mean	std
mfr						
General Mills	111.363636	10.371873	7.954545	3.872704	34.485852	8.946704
Kellogg's	108.695652	22.218818	7.565217	4.500768	44.038462	14.457434
Maltex	100.000000	NaN	3.000000	NaN	54.850917	NaN
Nabisco	86.666667	10.327956	1.833333	2.857738	67.968567	5.509326
Post	108.888889	10.540926	8.777778	4.576510	41.705744	10.047647
Quaker Oats	95.000000	29.277002	5.500000	4.780914	42.915990	16.797673
Ralston Purina	115.000000	22.677868	6.125000	3.563205	41.542997	6.080841

<string>:1: FutureWarning: ['name', 'type'] did not aggregate successfully. If any error is raised this will raise in a future version of pandas. Drop these columns/ops to avoid this warning.

Aggregation by column

```
1 cereal.groupby("mfr").agg({
2     "calories": ['min', 'max'],
3     "sugars":   ['mean', 'median'],
4     "rating":   ['sum', 'count']
5 })
```

	calories		sugars		rating	
	min	max	mean	median	sum	count
mfr						
General Mills	100	140	7.954545	8.5	758.688737	22
Kellogg's	50	160	7.565217	7.0	1012.884634	23
Maltex	100	100	3.000000	3.0	54.850917	1
Nabisco	70	100	1.833333	0.0	407.811403	6
Post	90	120	8.777778	10.0	375.351697	9
Quaker Oats	50	120	5.500000	6.0	343.327919	8
Ralston Purina	90	150	6.125000	5.5	332.343977	8

Named aggregation

It is also possible to use special syntax to aggregate specific columns into a named output column,

```
1 cereal.groupby("mfr", as_index=False).agg(  
2     min_cal = ("calories", "min"),  
3     max_cal = ("calories", "max"),  
4     med_sugar = ("sugars", "median"),  
5     avg_rating = ("rating", "mean")  
6 )
```

	mfr	min_cal	max_cal	med_sugar	avg_rating
0	General Mills	100	140	8.5	34.485852
1	Kellogg's	50	160	7.0	44.038462
2	Maltex	100	100	3.0	54.850917
3	Nabisco	70	100	0.0	67.968567
4	Post	90	120	10.0	41.705744
5	Quaker Oats	50	120	6.0	42.915990
6	Ralston Purina	90	150	5.5	41.542997

Transformation

The `transform()` method returns a DataFrame with the aggregated result matching the size (or length 1) of the input group(s),

```
1 cereal.groupby("mfr").transform(np.mean)
```

	calories	sugars	rating
0	86.666667	1.833333	67.968567
1	95.000000	5.500000	42.915990
2	108.695652	7.565217	44.038462
3	108.695652	7.565217	44.038462
4	115.000000	6.125000	41.542997
..
72	111.363636	7.954545	34.485852
73	111.363636	7.954545	34.485852
74	115.000000	6.125000	41.542997
75	111.363636	7.954545	34.485852
76	111.363636	7.954545	34.485852

[77 rows x 3 columns]

<string>:1: FutureWarning: The default value of numeric only in DataFrameGroupBy.mean is

```
1 cereal.groupby("type").transform("mean")
```

	calories	sugars	rating
0	107.162162	7.175676	42.095218
1	107.162162	7.175676	42.095218
2	107.162162	7.175676	42.095218
3	107.162162	7.175676	42.095218
4	107.162162	7.175676	42.095218
..
72	107.162162	7.175676	42.095218
73	107.162162	7.175676	42.095218
74	107.162162	7.175676	42.095218
75	107.162162	7.175676	42.095218
76	107.162162	7.175676	42.095218

[77 rows x 3 columns]

<string>:1: FutureWarning: The default value of numeric only in DataFrameGroupBy.mean is

Practical transformation

`transform()` will generally be most useful via a user defined function, the lambda argument is each column of each group.

```
1 ( cereal
2   .groupby("mfr")
3   .transform(
4     lambda x: (x - np.mean(x))/np.std(x)
5   )
6 )
```

	calories	sugars	rating
0	-1.767767	1.597191	0.086375
1	0.912871	0.559017	-0.568474
2	-1.780712	-0.582760	1.088220
3	-2.701081	-1.718649	3.512566
4	-0.235702	0.562544	-1.258442
..
72	-0.134568	-1.309457	0.528580
73	-0.134568	1.069190	-0.770226
74	-0.707107	-0.937573	1.449419
75	-1.121403	-1.309457	1.957022
76	-0.134568	0.012013	0.194681

[77 rows x 3 columns]

Filtering groups

`filter()` also respects groups and allows for the inclusion / exclusion of groups based on user specified criteria,

```
1 ( cereal
2   .groupby("mfr")
3   .filter(lambda x: len(x) > 8)
4 )
```

	name	mfr	...	sugars	rating
2	All-Bran	Kellogg's	...	5	59.425505
3	All-Bran with Extra Fiber	Kellogg's	...	0	93.704912
5	Apple Cinnamon Cheerios	General Mills	...	10	29.509541
6	Apple Jacks	Kellogg's	...	14	33.174094
7	Basic 4	General Mills	...	8	37.038562
9	Bran Flakes	Post	...	5	53.313813
11	Cheerios	General Mills	...	1	50.764999
12	Cinnamon Toast Crunch	General Mills	...	9	19.823573
13	Clusters	General Mills	...	7	40.400208
14	Cocoa Puffs	General Mills	...	13	22.736446
16	Corn Flakes	Kellogg's	...	2	45.863324
17	Corn Pops	Kellogg's	...	12	35.782791
18	Count Chocula	General Mills	...	13	22.396513
19	Cracklin' Oat Bran	Kellogg's	...	7	40.448772
21	Crispix	Kellogg's	...	3	46.895644
22	Crispy Wheat & Raisins	General Mills	...	10	36.176196

```

1 ( cereal
2   .groupby("mfr")
3   .size()
4 )

```

```

mfr
General Mills      22
Kellogg's          23
Maltex             1
Nabisco            6
Post              9
Quaker Oats       8
Ralston Purina    8
dtype: int64

```

```

1 ( cereal
2   .groupby("mfr")
3   .filter(lambda x: len(x) > 8)
4   .groupby("mfr")
5   .size()
6 )

```

```

mfr
General Mills      22
Kellogg's          23
Post              9
dtype: int64

```

