

# pandas

## Lecture 08

Dr. Colin Rundel

# pandas

pandas is an implementation of data frames in Python - it takes much of its inspiration from R and NumPy.

pandas aims to be the fundamental high-level building block for doing practical, real world data analysis in Python. Additionally, it has the broader goal of becoming the most powerful and flexible open source data analysis / manipulation tool available in any language.

Key features:

- DataFrame and Series (column) object classes
- Reading and writing tabular data
- Data munging (filtering, grouping, summarizing, joining, etc.)
- Data reshaping

# DataFrame

- Just like R a DataFrame is a collection of vectors with a common length
- Column types can be heterogeneous
- Both columns and rows can have names

```
1 iris = pd.read_csv("data/iris.csv")
2 type(iris)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
1 iris
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa
..	...	...	...	...	...
145	6.7	3.0	5.2	2.3	virginica
146	6.3	2.5	5.0	1.9	virginica
147	6.5	3.0	5.2	2.0	virginica
148	6.2	3.4	5.4	2.3	virginica
149	5.9	3.0	5.1	1.8	virginica

# Series

The columns of a DataFrame are constructed as Series - a 1d array like object containing values of the same type (similar to an ndarray).

```
1 pd.Series([1,2,3,4])
```

```
0    1
1    2
2    3
3    4
dtype: int64
```

```
1 pd.Series(["C","B","A"])
```

```
0    C
1    B
2    A
dtype: object
```

```
1 pd.Series([True])
```

```
0    True
dtype: bool
```

```
1 pd.Series(range(5))
```

```
0    0
1    1
2    2
3    3
4    4
dtype: int64
```

```
1 pd.Series([1,"A",True])
```

```
0    1
1    A
2    True
dtype: object
```

# Series methods

Once constructed the components of a series can be accessed via `array` and `index` attributes.

```
1 s = pd.Series([4,2,1,3])
```

```
1 s
```

```
0    4
1    2
2    1
3    3
dtype: int64
```

```
1 s.array
```

```
<PandasArray>
[4, 2, 1, 3]
Length: 4, dtype: int64
```

```
1 s.index
```

```
RangeIndex(start=0, stop=4, step=1)
```

An index can also be explicitly provided when constructing a Series,

```
1 t = pd.Series([4,2,1,3], index=["a","b","c","d"])
```

```
1 t
```

```
a    4
b    2
c    1
d    3
dtype: int64
```

```
1 t.array
```

```
<PandasArray>
[4, 2, 1, 3]
Length: 4, dtype: int64
```

```
1 t.index
```

```
Index(['a', 'b', 'c', 'd'], dtype='object')
```

# Series + NumPy

Series objects are compatible with NumPy like functions (vectorized)

```
1 t = pd.Series([4,2,1,3], index=["a","b","c","d"])
```

```
1 t + 1
```

```
a    5
b    3
c    2
d    4
dtype: int64
```

```
1 t / 2 + 1
```

```
a    3.0
b    2.0
c    1.5
d    2.5
dtype: float64
```

```
1 np.log(t)
```

```
a    1.386294
b    0.693147
c    0.000000
d    1.098612
dtype: float64
```

```
1 np.exp(-t**2/2)
```

```
a    0.000335
b    0.135335
c    0.606531
d    0.011109
dtype: float64
```

# Series indexing

Series can be indexed in the same way as NumPy arrays with the addition of being able to use index label(s) when selecting elements.

```
1 t = pd.Series([4,2,1,3], index=["a","b","c","d"])
```

```
1 t[1]
```

```
2
```

```
1 t[[1,2]]
```

```
b    2
c    1
dtype: int64
```

```
1 t["c"]
```

```
1
```

```
1 t[["a","d"]]
```

```
a    4
d    3
dtype: int64
```

```
1 t[t == 3]
```

```
d    3
dtype: int64
```

```
1 t[t % 2 == 0]
```

```
a    4
b    2
dtype: int64
```

```
1 t["d"] = 6
2 t
```

```
a    4
b    2
c    1
d    6
dtype: int64
```

# Index alignment

When performing (arithmetic) operations on series, they will attempt to align the operation by the index values,

```
1 m = pd.Series([1,2,3,4], index = ["a","b","c","d"])
2 n = pd.Series([4,3,2,1], index = ["d","c","b","a"])
3 o = pd.Series([1,1,1,1,1], index = ["b","d","a","c","e"])
```

```
1 m + n
```

```
a    2
b    4
c    6
d    8
dtype: int64
```

```
1 n + m
```

```
a    2
b    4
c    6
d    8
dtype: int64
```

```
1 n + o
```

```
a    2.0
b    3.0
c    4.0
d    5.0
e     NaN
dtype: float64
```



# Series and dicts

Series can also be constructed from dicts, in which case the keys are used to create the index,

```
1 d = {"anna": "A+", "bob": "B-", "carol": "C", "dave": "D+"}  
2 pd.Series(d)
```

```
anna      A+  
bob       B-  
carol     C  
dave      D+  
dtype: object
```

Index order will follow key order, unless overridden by [index](#),

```
1 pd.Series(d, index = ["dave", "carol", "bob", "anna"])
```

```
dave      D+  
carol     C  
bob       B-  
anna      A+  
dtype: object
```

# Missing values

Pandas encodes missing values using NaN (mostly),

```
1 s = pd.Series(  
2     {"anna": "A+", "bob": "B-",  
3     "carol": "C", "dave": "D+"},  
4     index = ["erin", "dave", "carol", "bob", "anna"]  
5 )
```

```
1 s
```

```
erin      NaN  
dave      D+  
carol      C  
bob       B-  
anna      A+  
dtype: object
```

```
1 pd.isna(s)
```

```
erin      True  
dave     False  
carol     False  
bob       False  
anna     False  
dtype: bool
```

```
1 s = pd.Series(  
2     {"anna": 97, "bob": 82,  
3     "carol": 75, "dave": 68},  
4     index = ["erin", "dave", "carol", "bob", "anna"],  
5     dtype = 'int64'  
6 )
```

```
1 s
```

```
erin      NaN  
dave     68.0  
carol     75.0  
bob       82.0  
anna     97.0  
dtype: float64
```

```
1 pd.isna(s)
```

```
erin      True  
dave     False  
carol     False  
bob       False  
anna     False  
dtype: bool
```

# Aside - why `np.isnan()`?

```
1 s = pd.Series([1,2,3,None])
2 s
```

```
0    1.0
1    2.0
2    3.0
3    NaN
dtype: float64
```

```
1 pd.isna(s)
```

```
0    False
1    False
2    False
3     True
dtype: bool
```

```
1 s == np.nan
```

```
0    False
1    False
2    False
3    False
dtype: bool
```

```
1 np.nan == np.nan
```

False

```
1 np.nan != np.nan
```

True

```
1 np.isnan(np.nan)
```

True

```
1 np.isnan(0)
```

False

# Native NAs

Recent versions of pandas have attempted to adopt a more native missing value, particularly for integer and boolean types,

```
1 pd.Series([1,2,3,None])
```

```
0    1.0
1    2.0
2    3.0
3    NaN
dtype: float64
```

```
1 pd.Series([True,False,None])
```

```
0    True
1   False
2    None
dtype: object
```

```
1 pd.isna( pd.Series([1,2,3,None]) )
```

```
0    False
1    False
2    False
3     True
dtype: bool
```

```
1 pd.isna( pd.Series([True,False,None]) )
```

```
0    False
1    False
2     True
dtype: bool
```

# Setting dtype

We can force things by setting the Series dtype,

```
1 pd.Series(  
2     [1,2,3,None],  
3     dtype = pd.Int64Dtype()  
4 )
```

```
0      1  
1      2  
2      3  
3  <NA>  
dtype: Int64
```

```
1 pd.Series(  
2     [True, False,None],  
3     dtype = pd.BooleanDtype()  
4 )
```

```
0      True  
1     False  
2     <NA>  
dtype: boolean
```

# String series

Series containing strings can be accessed via the `str` attribute,

```
1 s = pd.Series(["the quick", "brown fox", "jumps over", "a lazy dog"])
```

```
1 s
```

```
0    the quick
1    brown fox
2    jumps over
3    a lazy dog
dtype: object
```

```
1 s.str.upper()
```

```
0    THE QUICK
1    BROWN FOX
2    JUMPS OVER
3    A LAZY DOG
dtype: object
```

```
1 s.str.split(" ")
```

```
0    [the, quick]
1    [brown, fox]
2    [jumps, over]
3    [a, lazy, dog]
dtype: object
```

```
1 s.str.split(" ").str[1]
```

```
0    quick
1     fox
2     over
3     lazy
dtype: object
```

```
1 pd.Series([1,2,3]).str
```

Error: AttributeError: Can only use .str accessor with string values!

# Categorical Series

```
1 pd.Series(  
2     ["Mon", "Tue", "Wed", "Thur", "Fri"]  
3 )
```

```
0    Mon  
1    Tue  
2    Wed  
3    Thur  
4    Fri  
dtype: object
```

```
1 pd.Series(  
2     ["Mon", "Tue", "Wed", "Thur", "Fri"],  
3     dtype="category"  
4 )
```

```
0    Mon  
1    Tue  
2    Wed  
3    Thur  
4    Fri  
dtype: category  
Categories (5, object): ['Fri', 'Mon', 'Thur',  
                        'Tue', 'Wed']
```

```
1 pd.Series(  
2     ["Mon", "Tue", "Wed", "Thur", "Fri"],  
3     dtype=pd.CategoricalDtype(ordered=True)  
4 )
```

```
0    Mon  
1    Tue  
2    Wed  
3    Thur  
4    Fri  
dtype: category  
Categories (5, object): ['Fri' < 'Mon' < 'Thur' < 'Tue' < 'Wed']
```

# Category orders

```
1 pd.Series(  
2     ["Tue", "Thur", "Mon", "Sat"],  
3     dtype=pd.CategoricalDtype(categories=["Mon", "Tue", "Wed", "Thur", "Fri"],  
4 )
```

0 Tue

1 Thur

2 Mon

3 NaN

dtype: category

Categories (5, object): ['Mon' < 'Tue' < 'Wed' < 'Thur' < 'Fri']



# Constructing DataFrames

Earlier we saw reading a DataFrame in via `read_csv()`, but data frames can also be constructed via `DataFrame()`, in general this is done using a dictionary of columns:

```
1 n = 5
2 d = {
3     "id":      np.random.randint(100, 999, n),
4     "weight":  np.random.normal(70, 20, n),
5     "height":  np.random.normal(170, 15, n),
6     "date":    pd.date_range(start='2/1/2022',
7                               periods=n, freq='D')
8 }
9 d
```

```
1 df = pd.DataFrame(d)
2 df
```

	id	weight	height	date
0	623	54.731075	176.979498	2022-02-01
1	142	98.365587	175.065206	2022-02-02
2	336	60.018398	184.063187	2022-02-03
3	846	91.877850	171.272554	2022-02-04
4	570	33.720132	168.207309	2022-02-05

```
{'id': array([623, 142, 336, 846, 570]), 'weight':
array([54.73107509, 98.3655868 , 60.01839835,
91.87785014, 33.72013207]), 'height':
array([176.97949842, 175.06520576, 184.06318659,
171.27255432, 168.20730921]), 'date':
DatetimeIndex(['2022-02-01', '2022-02-02', '2022-
02-03', '2022-02-04',
               '2022-02-05'],
              dtype='datetime64[ns]', freq='D')}
```

# DataFrame from ndarray

For 2d ndarrays it is also possible to construct a DataFrame - generally it is a good idea to provide column names and row names (indexes)

```
1 pd.DataFrame(  
2     np.diag([1,2,3]),  
3     columns = ["x","y","z"]  
4 )
```

	x	y	z
0	1	0	0
1	0	2	0
2	0	0	3

```
1 pd.DataFrame(  
2     np.diag([1,2,3]),  
3     columns = ["x","y","z"]  
4 )
```

	x	y	z
0	1	0	0
1	0	2	0
2	0	0	3

```
1 pd.DataFrame(  
2     np.tri(5,3,-1),  
3     columns = ["x","y","z"],  
4     index = ["a","b","c","d","e"]  
5 )
```

	x	y	z
a	0.0	0.0	0.0
b	1.0	0.0	0.0
c	1.0	1.0	0.0
d	1.0	1.0	1.0
e	1.0	1.0	1.0

# DataFrame indexing

```
1 df
```

	id	weight	height	date
0	623	54.731075	176.979498	2022-02-01
1	142	98.365587	175.065206	2022-02-02
2	336	60.018398	184.063187	2022-02-03
3	846	91.877850	171.272554	2022-02-04
4	570	33.720132	168.207309	2022-02-05

Selecting a column,

```
1 df[0]
```

Error: KeyError: 0

```
1 df["id"]
```

0	623
1	142
2	336
3	846
4	570

Name: id, dtype: int64

```
1 df.id
```

0	623
1	142
2	336
3	846
4	570

Name: id, dtype: int64

Selecting rows (a single slice is assumed to refer to the rows)

```
1 df[1:3]
```

	id	weight	height	date
1	142	98.365587	175.065206	2022-02-02
2	336	60.018398	184.063187	2022-02-03

```
1 df[0::2]
```

	id	weight	height	date
0	623	54.731075	176.979498	2022-02-01
2	336	60.018398	184.063187	2022-02-03
4	570	33.720132	168.207309	2022-02-05

# Index by position

```
1 df
```

	id	weight	height	date
0	623	54.731075	176.979498	2022-02-01
1	142	98.365587	175.065206	2022-02-02
2	336	60.018398	184.063187	2022-02-03
3	846	91.877850	171.272554	2022-02-04
4	570	33.720132	168.207309	2022-02-05

```
1 df.iloc[1]
```

	id	weight	height	date
1	142	98.365587	175.065206	2022-02-02 00:00:00

Name: 1, dtype: object

```
1 df.iloc[[1]]
```

	id	weight	height	date
1	142	98.365587	175.065206	2022-02-02

```
1 df.iloc[0:2]
```

	id	weight	height	date
0	623	54.731075	176.979498	2022-02-01
1	142	98.365587	175.065206	2022-02-02

```
1 df.iloc[lambda x: x.index % 2 != 0]
```

	id	weight	height	date
1	142	98.365587	175.065206	2022-02-02
3	846	91.877850	171.272554	2022-02-04

```
1 df.iloc[1:3,1:3]
```

	weight	height
1	98.365587	175.065206
2	60.018398	184.063187

```
1 df.iloc[0:3, [0,3]]
```

	id	date
0	623	2022-02-01
1	142	2022-02-02
2	336	2022-02-03

```
1 df.iloc[0:3, [True, True, False, False]]
```

	id	weight
0	623	54.731075
1	142	98.365587
2	336	60.018398

# Index by name

```
1 df.index = (["anna","bob","carol", "dave", "erin"])
2 df
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01
bob	142	98.365587	175.065206	2022-02-02
carol	336	60.018398	184.063187	2022-02-03
dave	846	91.877850	171.272554	2022-02-04
erin	570	33.720132	168.207309	2022-02-05

```
1 df.loc["anna"]
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01 00:00:00

Name: anna, dtype: object

```
1 df.loc[["anna"]]
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01

```
1 df.loc["bob":"dave"]
```

	id	weight	height	date
bob	142	98.365587	175.065206	2022-02-02
carol	336	60.018398	184.063187	2022-02-03
dave	846	91.877850	171.272554	2022-02-04

```
1 df.loc[df.id < 300]
```

	id	weight	height	date
bob	142	98.365587	175.065206	2022-02-02

```
1 df.loc[:, "date"]
```

	date
anna	2022-02-01
bob	2022-02-02
carol	2022-02-03
dave	2022-02-04
erin	2022-02-05

Name: date, dtype: datetime64[ns]

```
1 df.loc[["bob","erin"], "weight":"height"]
```

	weight	height
bob	98.365587	175.065206
erin	33.720132	168.207309

```
1 df.loc[0:2, "weight":"height"]
```

Error: TypeError: cannot do slice indexing on Index with these indexers [0] of type int

# Views vs. Copies

In general most pandas operations will generate a new object but some will return views, mostly the later occurs with subsetting.

```
1 d = pd.DataFrame(np.arange(6).reshape(3,2), columns=['x', 'y'])
2 d
```

	x	y
0	0	1
1	2	3
2	4	5

```
1 v = d.iloc[0:2,0:2]
2 v
```

	x	y
0	0	1
1	2	3

```
1 d.iloc[0,1] = -1
2 v
```

	x	y
0	0	-1
1	2	3

```
1 v.iloc[0,0] = np.pi
2 v
```

	x	y
0	3.141593	-1
1	2.000000	3

```
1 d
```

	x	y
0	0	-1
1	2	3
2	4	5

# Filtering rows

The `query()` method can be used for filtering rows, it evaluates a string expression in the context of the data frame.

```
1 df
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01
bob	142	98.365587	175.065206	2022-02-02
carol	336	60.018398	184.063187	2022-02-03
dave	846	91.877850	171.272554	2022-02-04
erin	570	33.720132	168.207309	2022-02-05

```
1 df.query('date == "2022-02-01"')
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01

```
1 df.query('weight > 50')
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01
bob	142	98.365587	175.065206	2022-02-02
carol	336	60.018398	184.063187	2022-02-03
dave	846	91.877850	171.272554	2022-02-04

```
1 df.query('weight > 50 & height < 165')
```

Empty DataFrame  
Columns: [id, weight, height, date]  
Index: []

```
1 qid = 414  
2 df.query('id == @qid')
```

Empty DataFrame  
Columns: [id, weight, height, date]  
Index: []

# Element access

```
1 df
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01
bob	142	98.365587	175.065206	2022-02-02
carol	336	60.018398	184.063187	2022-02-03
dave	846	91.877850	171.272554	2022-02-04
erin	570	33.720132	168.207309	2022-02-05

```
1 df[0,0]
```

Error: KeyError: (0, 0)

```
1 df.iat[0,0]
```

623

```
1 df.id[0]
```

623

```
1 df[0:1].id[0]
```

623

```
1 df["anna", "id"]
```

Error: KeyError: ('anna', 'id')

```
1 df.at["anna", "id"]
```

623

```
1 df["id"]["anna"]
```

623

```
1 df["id"][0]
```

623



# DataFrame properties

```
1 df.size
```

```
20
```

```
1 df.shape
```

```
(5, 4)
```

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 5 entries, anna to erin
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0   id       5 non-null       int64
1   weight   5 non-null       float64
2   height   5 non-null       float64
3   date     5 non-null       datetime64[ns]
dtypes: datetime64[ns](1), float64(2), int64(1)
memory usage: 372.0+ bytes
```

```
1 df.dtypes
```

```
id           int64
weight       float64
height       float64
date         datetime64[ns]
dtype: object
```

```
1 df.describe()
```

	id	weight	height
count	5.000000	5.000000	5.000000
mean	503.400000	67.742608	175.117551
std	271.453127	26.957230	6.042132
min	142.000000	33.720132	168.207309
25%	336.000000	54.731075	171.272554
50%	570.000000	60.018398	175.065206
75%	623.000000	91.877850	176.979498
max	846.000000	98.365587	184.063187

# Selecting Columns

Beyond the use of `loc()` and `iloc()` there is also the `filter()` method which can be used to select columns (or indices) by name with pattern matching

```
1 df.filter(items=["id","weight"])
```

	id	weight
anna	623	54.731075
bob	142	98.365587
carol	336	60.018398
dave	846	91.877850
erin	570	33.720132

```
1 df.filter(like = "i")
```

	id	weight	height
anna	623	54.731075	176.979498
bob	142	98.365587	175.065206
carol	336	60.018398	184.063187
dave	846	91.877850	171.272554
erin	570	33.720132	168.207309

```
1 df.filter(regex="ght$")
```

	weight	height
anna	54.731075	176.979498
bob	98.365587	175.065206
carol	60.018398	184.063187
dave	91.877850	171.272554
erin	33.720132	168.207309

```
1 df.filter(like="o", axis=0)
```

	id	weight	height	date
bob	142	98.365587	175.065206	2022-02-02
carol	336	60.018398	184.063187	2022-02-03

# Adding columns

Indexing with assignment allows for inplace modification of a DataFrame, while `assign()` creates a new object (but is chainable)

```
1 df['student'] = [True, True, True, False, None]
2 df['age'] = [19, 22, 25, None, None]
3 df
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
dave	846	91.877850	171.272554	2022-02-04	False	NaN
erin	570	33.720132	168.207309	2022-02-05	None	NaN

```
1 df.assign(
2     student = lambda x: np.where(x.student, "yes", "no"),
3     rand = np.random.rand(5)
4 )
```

	id	weight	height	date	student	age	rand
anna	623	54.731075	176.979498	2022-02-01	yes	19.0	0.752240
bob	142	98.365587	175.065206	2022-02-02	yes	22.0	0.876159
carol	336	60.018398	184.063187	2022-02-03	yes	25.0	0.797148
dave	846	91.877850	171.272554	2022-02-04	no	NaN	0.863391
erin	570	33.720132	168.207309	2022-02-05	no	NaN	0.527163

```
1 df
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
dave	846	91.877850	171.272554	2022-02-04	False	NaN
erin	570	33.720132	168.207309	2022-02-05	None	NaN

# Removing columns (and rows)

Columns can be dropped via the `drop()` method,

```
1 df.drop(['student'])
```

Error: KeyError: "['student'] not found in axis"

```
1 df.drop(['student'], axis=1)
```

	id	weight	height	date	age
anna	623	54.731075	176.979498	2022-02-01	19.0
bob	142	98.365587	175.065206	2022-02-02	22.0
carol	336	60.018398	184.063187	2022-02-03	25.0
dave	846	91.877850	171.272554	2022-02-04	NaN
erin	570	33.720132	168.207309	2022-02-05	NaN

```
1 df.drop(['anna', 'dave'])
```

	id	weight	height	date	student	age
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
erin	570	33.720132	168.207309	2022-02-05	None	NaN

```
1 df.drop(columns = df.columns == "age")
```

Error: KeyError: '[False, False, False, False, False, True] not found in axis'

```
1 df.drop(columns = df.columns[df.columns == "age"])
```

	id	weight	height	date	student
anna	623	54.731075	176.979498	2022-02-01	True
bob	142	98.365587	175.065206	2022-02-02	True
carol	336	60.018398	184.063187	2022-02-03	True
dave	846	91.877850	171.272554	2022-02-04	False
erin	570	33.720132	168.207309	2022-02-05	None

```
1 df.drop(columns = df.columns[df.columns.str.contains("ght")])
```

	id	date	student	age
anna	623	2022-02-01	True	19.0
bob	142	2022-02-02	True	22.0
carol	336	2022-02-03	True	25.0
dave	846	2022-02-04	False	NaN
erin	570	2022-02-05	None	NaN

# Dropping missing values

Columns can be dropped via the `drop()` method,

```
1 df
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
dave	846	91.877850	171.272554	2022-02-04	False	NaN
erin	570	33.720132	168.207309	2022-02-05	None	NaN

```
1 df.dropna()
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0

```
1 df.dropna(how="all")
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
dave	846	91.877850	171.272554	2022-02-04	False	NaN
erin	570	33.720132	168.207309	2022-02-05	None	NaN

```
1 df.dropna(axis=1)
```

	id	weight	height	date
anna	623	54.731075	176.979498	2022-02-01
bob	142	98.365587	175.065206	2022-02-02
carol	336	60.018398	184.063187	2022-02-03
dave	846	91.877850	171.272554	2022-02-04
erin	570	33.720132	168.207309	2022-02-05

```
1 df.dropna(axis=1, thresh=4)
```

	id	weight	height	date	student
anna	623	54.731075	176.979498	2022-02-01	True
bob	142	98.365587	175.065206	2022-02-02	True
carol	336	60.018398	184.063187	2022-02-03	True
dave	846	91.877850	171.272554	2022-02-04	False
erin	570	33.720132	168.207309	2022-02-05	None

# Sorting

DataFrames can be sorted on one or more columns via `sort_values()`,

```
1 df
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
dave	846	91.877850	171.272554	2022-02-04	False	NaN
erin	570	33.720132	168.207309	2022-02-05	None	NaN

```
1 df.sort_values(by=["student","id"], ascending=[True,False])
```

	id	weight	height	date	student	age
dave	846	91.877850	171.272554	2022-02-04	False	NaN
anna	623	54.731075	176.979498	2022-02-01	True	19.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
erin	570	33.720132	168.207309	2022-02-05	None	NaN

# Row binds

DataFrames can have their rows joined via the the `concat()` function (`append()` is also available but deprecated),

```
1 df1 = pd.DataFrame(  
2     np.arange(6).reshape(3,2),  
3     columns=list("xy")  
4 )  
5 df1
```

	x	y
0	0	1
1	2	3
2	4	5

```
1 pd.concat([df1,df2])
```

	x	y
0	0	1
1	2	3
2	4	5
0	12	11
1	10	9
2	8	7

```
1 df2 = pd.DataFrame(  
2     np.arange(12,6,-1).reshape(3,2),  
3     columns=list("xy")  
4 )  
5 df2
```

	x	y
0	12	11
1	10	9
2	8	7

```
1 pd.concat([df1.loc[:,["y","x"]],df2])
```

	y	x
0	1	0
1	3	2
2	5	4
0	11	12
1	9	10
2	7	8

# Imputing columns

When binding rows missing columns will be added with NaN or <NA> entries.

```
1 df3 = pd.DataFrame(np.ones((3,3)), columns=list("xbz"))
2 df3
```

	x	b	z
0	1.0	1.0	1.0
1	1.0	1.0	1.0
2	1.0	1.0	1.0

```
1 pd.concat([df1,df3,df2])
```

	x	y	b	z
0	0.0	1.0	NaN	NaN
1	2.0	3.0	NaN	NaN
2	4.0	5.0	NaN	NaN
0	1.0	NaN	1.0	1.0
1	1.0	NaN	1.0	1.0
2	1.0	NaN	1.0	1.0
0	12.0	11.0	NaN	NaN
1	10.0	9.0	NaN	NaN
2	8.0	7.0	NaN	NaN



# Column binds

Similarly, columns can be joined with `concat()` where `axis=1`,

```
1 df1 = pd.DataFrame(  
2     np.arange(6).reshape(3,2),  
3     columns=list("xy"),  
4     index=list("abc")  
5 )  
6 df1
```

	x	y
a	0	1
b	2	3
c	4	5

```
1 pd.concat([df1,df2], axis=1)
```

	x	y	m	n
a	0	1	10.0	9.0
b	2	3	NaN	NaN
c	4	5	8.0	7.0

```
1 df2 = pd.DataFrame(  
2     np.arange(10,6,-1).reshape(2,2),  
3     columns=list("mn"),  
4     index=list("ac")  
5 )  
6 df2
```

	m	n
a	10	9
c	8	7

```
1 pd.concat([df1,df2], axis=1, join="inner")
```

	x	y	m	n
a	0	1	10	9
c	4	5	8	7

# Joining DataFrames

Table joins are implemented via the `merge()` function or method,

```
1 df1 = pd.DataFrame(  
2     {'a': ['foo', 'bar'], 'b': [1, 2]}  
3 )  
4 df1
```

	a	b
0	foo	1
1	bar	2

```
1 pd.merge(df1,df2, how="inner")
```

	a	b	c
0	foo	1	3

```
1 pd.merge(df1,df2, how="outer", on="a")
```

	a	b	c
0	foo	1.0	3.0
1	bar	2.0	NaN
2	baz	NaN	4.0

```
1 df2 = pd.DataFrame(  
2     {'a': ['foo', 'baz'], 'c': [3, 4]}  
3 )  
4 df2
```

	a	c
0	foo	3
1	baz	4

```
1 df1.merge(df2, how="left")
```

	a	b	c
0	foo	1	3.0
1	bar	2	NaN

```
1 df1.merge(df2, how="right")
```

	a	b	c
0	foo	1.0	3
1	baz	NaN	4

# join vs merge vs concat

All three can be used to accomplish the same thing, in terms of “column bind” type operations.

- `concat()` stacks DataFrames on either axis, with basic alignment based on (row) indexes. `join` argument only supports “inner” and “outer”.
- `merge()` aligns based on one or more shared columns. `how` supports “inner”, “outer”, “left”, “right”, and “cross”.
- `join()` uses `merge()` behind the scenes, but prefers to join based on (row) indexes. Also has different default `how` compared to `merge()`, “left” vs “inner”.

# groupby and agg

Groups can be created within a DataFrame via `groupby()` - these groups are then used by the standard summary methods (e.g. `sum()`, `mean()`, `std()`, etc.).

```
1 df.groupby("student")
```

```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x2921ee500>
```

```
1 df.groupby("student").groups
```

```
{False: ['dave'], True: ['anna', 'bob', 'carol']}
```

```
1 df.groupby("student").mean(numeric_only=True)
```

	id	weight	height	age
student				
False	846.0	91.877850	171.272554	NaN
True	367.0	71.038353	178.702630	22.0

```
1 df.groupby("student", dropna=False).groups
```

```
Error: ValueError: Categorical categories cannot be null
```

```
1 df.groupby("student", dropna=False).mean(numeric_onl
```

	id	weight	height	age
student				
False	846.0	91.877850	171.272554	NaN
True	367.0	71.038353	178.702630	22.0
NaN	570.0	33.720132	168.207309	NaN

# Selecting groups

```
1 df
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0
dave	846	91.877850	171.272554	2022-02-04	False	NaN
erin	570	33.720132	168.207309	2022-02-05	None	NaN

```
1 df.groupby("student").get_group(True)
```

	id	weight	height	date	student	age
anna	623	54.731075	176.979498	2022-02-01	True	19.0
bob	142	98.365587	175.065206	2022-02-02	True	22.0
carol	336	60.018398	184.063187	2022-02-03	True	25.0

```
1 df.groupby("student").get_group(False)
```

	id	weight	height	date	student	age
dave	846	91.87785	171.272554	2022-02-04	False	NaN

```
1 df.groupby("student", dropna=False).get_group(np.nan)
```

Error: KeyError: nan

# Aggregation

```
1 df = df.drop("date", axis=1)
```

```
1 df.groupby("student").agg("mean")
```

	id	weight	height	age
student				
False	846.0	91.877850	171.272554	NaN
True	367.0	71.038353	178.702630	22.0

```
1 df.groupby("student").agg([np.mean, np.std])
```

	id		weight	...	height	age	
	mean	std	mean	...	std	mean	std
student				...			
False	846.0	NaN	91.877850	...	NaN	NaN	NaN
True	367.0	241.993802	71.038353	...	4.740021	22.0	3.0

[2 rows x 8 columns]

